
System Memory Management Unit (SMMU) Overview



Qualcomm Technologies, Inc.

80-P2055-1 B

Confidential and Proprietary – Qualcomm Technologies, Inc.

Restricted Distribution: Not to be distributed to anyone who is not an employee of either Qualcomm Technologies, Inc. or its affiliated companies without the express approval of Qualcomm Configuration Management.



Confidential and Proprietary – Qualcomm Technologies, Inc.

Qualcomm
Confidential - May Contain Trade Secrets
2024-03-15 06:16:34 GMT
rajkumar.veer@harman.com

NO PUBLIC DISCLOSURE PERMITTED: Please report postings of this document on public servers or websites to: DocCtrlAgent@qualcomm.com.

Not to be used, copied, reproduced, or modified in whole or in part, nor its contents revealed in any manner to others without the express written permission of Qualcomm Technologies, Inc.

Qualcomm is a trademark of Qualcomm Incorporated, registered in the United States and other countries. All Qualcomm Incorporated trademarks are used with permission. Other product and brand names may be trademarks or registered trademarks of their respective owners.

This technical data may be subject to U.S. and international export, re-export, or transfer ("export") laws. Diversion contrary to U.S. and international law is strictly prohibited.

Qualcomm Technologies, Inc.
5775 Morehouse Drive
San Diego, CA 92121
U.S.A.

© 2015 Qualcomm Technologies, Inc. and/or its affiliated companies. All rights reserved.

Revision History

Revision	Date	Description
A	November 2015	Initial release
B	December 2015	Slide 19 was added

Qualcomm
Confidential - May Contain Trade Secrets
2024-03-15 06:16:34 GMT
rajkumar.veer@harman.com

Contents

- SMMU Overview
- Key Terminology
- Important APIs
- Device Tree Setup
- API Usage
- SMMU Faults
- Important SMMU registers
- Fault Status Register
- Other registers (Important bits)
- SMMU Faults Examples
- References
- Questions?

SMMU Overview

- SMUU is basically an MMU for devices
- SMMU hardware block allows virtually contiguous memory to be backed by physically noncontiguous pages
- Memory translation logic in the SMMU is the same as the logic in the CPU MMU
- SMMU and IOMMU are used interchangeably

Key Terminology

- **Context Bank**
 - Set of registers that defines a translation context; every context bank in the system has the same set of registers
- **Stream ID**
 - Identifier used as input to the SMMU to uniquely identify the current transaction stream
- **IOVA**
 - I/O Virtual Address, a.k.a DMA address and Virtual address
- **Page Tables**
 - Tables in memory that define the virtual-to-physical mappings; SMMU walks these tables to perform translations
- **TTBR0**
 - Register holding the base address of the page tables

Key Terminology (cont.)

- Page Level
 - Level of the page table walk; generally only comes up when looking at translation faults
- Translation Stage
 - Stage of the translation process; Stage 1 is managed by HLOS; Stage 2 is managed by the hypervisor
- IPA
 - Intermediate Physical address; output address from Stage 1 and input address to Stage 2
- TLB
 - Translation Lookaside Buffer; cached IOVA to Phys mapping
- Bypass
 - When a stage (stage 1, stage 2, or both) is not considered for translation (i.e., input = output)

Important APIs



API	Description
arm_iommu_create_mapping	configure a VA range
arm_iommu_attach_device	associate a VA range with a device and enable SMMU translations
dma_alloc_coherent	allocate buffer and map into SMMU for device
dma_map_sg	map existing buffer (by struct scatterlist) into SMMU for device
ion_share_dma_buf	Convert an Ion handle to a dma buf
dma_buf_attach	Prepare a dma buf for dma on the given device
dma_buf_map_attachment	Get the scatterlist for the given attachment
msm_iommu_get_bus	Get proper bus for your domain

Device Tree Setup

- Top level device tree nodes for SMMUs. Clients will use phandles to this node in their config

File: [arch/arm64/boot/dts/qcom/msm-arm-smmu.dtsi](#)

```
jpeg_smmu: arm,smmu-jpeg@fda64000 {  
    compatible = "qcom,smmu-v2";  
    /* ... */  
};
```

- Client node (for a device driver that only needs one context bank)

Example:

```
qcom,my_device@fda1c000 {  
    compatible = "qcom,my-driver";  
    /* ...your usual driver code here... */  
    iommus = <&jpeg_smmu 0>,  
            <&jpeg_smmu 1>;  
};
```

- With this DT setup, device pointer for qcom, my_device will be used everywhere that a device pointer is required for DMA APIs.

Device Tree Setup (cont.)

- Client node (for a device driver that needs multiple context banks)

```
qcom,my_device@fda1c000 {  
    compatible = "qcom,my-driver";  
    /* ...your usual driver code here... */  
  
    /* context bank "sub-devices": */  
    iommu_test_device_cb1 {  
        compatible = "qcom,my-driver-new-cb";  
        iommus = <&jpeg_smmu 0>,  
                <&jpeg_smmu 1>;  
    };  
  
    iommu_test_device_cb2 {  
        compatible = "qcom,my-driver-new-cb";  
        iommus = <&jpeg_smmu 2>;  
    };  
};
```

- With this setup, the device pointers associated with the "sub-device" DT nodes will be used for the various DMA APIs.

API Usage: Mapping a buffer into an SMMU with DMA APIs

- Define the virtual address range to be used for allocations

Example:

```
struct dma_iommu_mapping *mapping;  
dma_addr_t va_start = 0;  
size_t va_len = 0x40000000;  
struct bus_type *bus;
```

```
bus = msm_iommu_get_bus(dev); // Pass the context bank device pointer here  
as discussed on the previous slides  
mapping = arm_iommu_create_mapping(bus, va_start, va_len);
```

This allocates memory for the page table

- After creating the structure `dma_iommu_mapping`, associate the mapping with the device:

```
arm_iommu_attach_device(cb_device, mapping); // context bank device pointer  
and mapping to be passed
```

API Usage: Allocating and Mapping a New Buffer

- Allocate a buffer which will be mapped into the SMMU:

Example:

```
size_t size = 0x1000;
void *cpu_addr;
dma_addr_t iova;
cpu_addr = dma_alloc_coherent(cb_device, size, &iova, GFP_KERNEL);
```

where, cb_device is your context bank device pointer

- dma_alloc_coherent() automatically maps the buffer into the SMMU if arm_iommu_attach_device() is called for this device before calling dma_alloc_coherent().

API Usage: Mapping an Existing Buffer

- If there is an existing buffer that was allocated in some other way besides `dma_alloc_*` (e.g. Ion memory, `kmalloc`), one can map the buffer into the SMMU with the `dma_map_*` APIs

- `struct dma_buf*` is needed

```
struct dma_buf *buf = ion_share_dma_buf(client, handle);           // if we have a
struct ion_handle
OR
struct dma_buf *buf = dma_buf_get(fd); // through a dma_buf file descriptor
```

- Once there is a `struct dma_buf*` available, map it into the SMMU as follows:

```
struct dma_buf_attachment *attach = dma_buf_attach(buf, cb_device);
struct sg_table *table = dma_buf_map_attachment(attach, DMA_TO_DEVICE);
```

```
dma_map_sg(cb_device, table->sgl, table->nents, DMA_TO_DEVICE);
where cb_device is the context bank device pointer.
```

- Supported values for the direction argument to `dma_map_sg()` are mapped to IOMMU page protection attributes as follows:

- | | |
|----------------------------------|---|
| ▪ <code>DMA_TO_DEVICE</code> | Device can read from the memory |
| ▪ <code>DMA_FROM_DEVICE</code> | Device can write to the memory |
| ▪ <code>DMA_BIDIRECTIONAL</code> | Device can read from or write to the memory |

API Usage: Mapping an Existing Buffer (cont.)

- After mapping, the output address (which is given to the device) can be obtained with `sg_dma_address(table->sgl)` and the length can be obtained with `sg_dma_len(table->sgl)`.

- After one is finished with the buffer, it must be unmapped

```
dma_unmap_sg(cb_device, table->sgl, table->nents, DMA_TO_DEVICE);  
dma_buf_unmap_attachment(attach, table, DMA_TO_DEVICE);  
dma_buf_detach(dma_buf, attach);  
dma_buf_put(dma_buf);
```

▪ **Clocks and Power**

- SMMU clients need to control clocks and power on MSM8996
 - SMMU Clients **do not** need to enable clocks and power when they call into the SMMU driver (through the DMA APIs) to set up mappings.
 - SMMU Clients **do** need to enable clocks and power when they expect the SMMU to be servicing transactions.

SMMU Faults

- Two types of SMMU Faults
 - Global Fault
 - Context Fault
- Global Fault
 - Occurs when either:
 - SMMU translation, when being processed, has no associated translation context bankor
 - A translation context bank is not the appropriate place to record the fault.
 - Fatal by default
 - Shows up in TZ/Hypervisor logs
 - Examples:
 - Unidentified stream faults: no match for the given stream ID was found in the stream matching table
 - Configuration Access Faults: Missing clocks, power.

SMMU Faults (cont.)

- Context Fault
 - Associated with a particular translation context bank
 - Categorized as Stage-1 and Stage-2 context faults
 - Stage-1 context faults are non-fatal and show up in the kernel log
 - Stage-2 context faults are fatal and show up in the hypervisor log
 - Examples
 - Translation Fault: An unmapped address was accessed
 - Permission Fault: Attempt to write to read-only memory, fetch instructions from non-executable memory, etc.

Important SMMU registers

- **FSR**
 - Fault Status Register; identifies the type of fault (translation, permission etc.)
- **FAR**
 - Fault Address Register; provides the input address that caused the fault
- **FSYNRm**
 - Fault Syndrome Registers; provides more information about the faulting access
- **CBFRSYNRA**
 - Context Bank Fault Restricted Syndrome Register A; provides the SID
- **TTBRm**
 - TTBR0 holds the base address of translation table 0
 - TTBR1 holds the base address of translation table 1

Fault Status Register

▪ Context Faults

- Bit 0: Reserved
- Bit 1: Translation Fault
- Bit 2: Access Fault
- Bit 3: Permission Fault
- Bit 4: External fault
- Bit 5: Match conflict fault
- Bit 6: TLB Lock Fault
- Bit 8: Address Size Fault
- Bit 9: Unsupported upstream transaction

▪ Global Faults

- Bit 0: Invalid context fault
- Bit 1: Unidentified Stream ID Fault
- Bit 2: Stream Match Conflict Fault
- Bit 3: Unimplemented Context Bank Fault
- Bit 4: Unimplemented Context Interrupt
- Bit 5: Configuration Access Fault
- Bit 6: External Fault
- Bit 7: Permission Fault
- Bit 8: Unsupported Upstream Transaction

Other Registers (Important Bits)

- Context Bank Fault Restricted Syndrome Register (CBFRSYNRn)
 - Gives fault syndrome information about the access that caused an exception in the associated translation context bank.
 - StreamID- Bits[15:0]
 - StreamID of the transaction that caused the fault
- Context Bank Fault Syndrome Register (CBn_FSYNRn)
 - Holds fault syndrome information about the memory access that caused a synchronous abort exception
 - Bit[4] = WNR -- Write Not Read. The possible values of this bit are:
 - 0 Read.
 - 1 Write

SMMU Faults Examples

A) Context Fault from Venus (TZ log)

SMMU:>> VENUS CB8 Fault:

FSR=0x40000402

>> Translation Fault

FAR=0x000000000004e0000

>> Faulting address = 0x4e0000

IPAFAR=0x000000000004e0000

FSYNR0=0x00080023

FSYNR1=0x170c0000

CBFRSYNRA8=0x00010180

>> SID = 0x180 (lower 16 bits)

.....

B) Context Fault from Display (TZ log)

SMMU:>> MDP CB2 Fault:

FSR=0x40000408

>> Permission Fault

FAR=0x00000000048a1000

>> Faulting address = 0x48a1000

IPAFAR=0x0000000000000000

FSYNR0=0x00000027

FSYNR1=0x1c03000a

CBFRSYNRA2=0x00000000

.....

C) Global Fault from Venus (TZ log)

SMMU GLOBAL CLIENT NON-SEC FAULT: bit mask=0x00000002

SMMU:>> VENUS NonSec Global Fault:

NSGFSR=0x80000002

>> FSR : Unidentified Stream ID Fault

NSGFAR=0x00000000deadd000

NSGFSYNR0=0x0000000c

NSGFSYNR1=0x00000006

NSGFSYNR2=0x17000000

NSCR0=0x00201e36

References

Documents	
Resources	
ARM System Memory Management Unit Architecture Specification	http://infocenter.arm.com/help/topic/com.arm.doc.ihi0062d.b/IHI0062D_b_system_mmu_architecture_specification.pdf

Qualcomm
Confidential - May Contain Trade Secrets
2024-03-15 06:16:34 GMT
rajkumar.veer@harman.com

Qualcomm
Confidential - May Contain Trade Secrets
2024-03-15 06:16:34 GMT
rajkumar.veer@harman.com

Questions?

<https://support.cdmatech.com>

